

Lecture 2: Image Classification

COMP 4970 / 7970

Deep Learning

Register for Free cloud compute credits

1. Amazon AWS Educate : \$100 credit / student

- Refs:

- <https://aws.amazon.com/education/awseducate/>



2. Google Cloud Platform: \$350 credit / student

- Free \$300 credit upon registration
- Additional \$50 from this class (thanks to Google!)



- Refs:

- <http://cs231n.github.io/gce-tutorial/>
 - <https://cloud.google.com/getting-started/>

- Necessary for Projects (if you use GPUs)

- Optional for Assignments

Administrative

- Use Piazza for questions/discussions on
 - Sign-up: <http://piazza.com/auburn/spring2018/comp49707970>
 - Assignments
 - Presentation topic and papers
 - Proposal / Projects
 - Anything regarding Deep Learning
- Presentation Topic + Paper
 - Schedule: <http://s.anhnguyen.me/2018/7970/schedule.htm>
 - 1. Choose a topic
 - 2. Choose a paper
 - Tips:
 - You are excited about!
 - Has talk/presentation videos online
 - Has demo
 - Has open-source code



ion	1/29	IT	Deep Networks, part III	
	2/2	IT	Deep Networks with Stochastic Depth (Huang et al 2016)	Grou
	2/5	IT	Introduction to Caffe	

Paper here only suggestion

- <http://cs231n.github.io/assignments2017/assignment1/>
- Deliverable: Zip up your code as `A1_<id>_<firstname>_<lastname>.zip` and upload to Canvas.

In this assignment you will practice putting together a simple image classification pipeline, based on the k-Nearest Neighbor or the SVM/Softmax classifier. The goals of this assignment are as follows:

- understand the basic **Image Classification pipeline** and the data-driven approach (train/predict stages)
- understand the train/val/test **splits** and the use of validation data for **hyperparameter tuning**.
- develop proficiency in writing efficient **vectorized** code with numpy
- implement and apply a k-Nearest Neighbor (**kNN**) classifier
- implement and apply a Multiclass Support Vector Machine (**SVM**) classifier
- implement and apply a **Softmax** classifier
- implement and apply a **Two layer neural network** classifier
- understand the differences and tradeoffs between these classifiers
- get a basic understanding of performance improvements from using **higher-level representations** than raw pixels (e.g. color histograms, Histogram of Gradient (HOG) features)

Project ideas

>350 Datasets Available at:
<http://archive.ics.uci.edu/ml/>

Newest Data Sets:		Most Popular Data Sets (hits since 2007):	
2012-07-17:	 Skin Segmentation	352287:	 Iris
2012-07-17:	 Planning Relax	249987:	 Adult
2012-07-04:	 Nomao	219203:	 Wine
2012-06-22:	 SMS Spam Collection	179585:	 Breast Cancer Wisconsin (Diagnostic)
2012-06-09:	 OPPORTUNITY Activity Recognition	165076:	 Car Evaluation
2012-05-21:	 ILPD (Indian Liver Patient Dataset)	139614:	 Abalone
2012-04-25:	 Gas Sensor Array Drift Dataset	125301:	 Poker Hand
	 Multivariate Classification		

Project ideas

Image/video database categories:

- Action Databases
- Attribute recognition
- Autonomous Driving
- Biological/Medical
- Camera calibration
- Face and Eye/Iris Databases
- Fingerprints
- General Images
- General RGBD and depth datasets
- General Videos
- Hand, Hand Grasp, Hand Action and Gesture Databases
- Image, Video and Shape Database Retrieval
- Object Databases
- People (static), human body pose
- People Detection and Tracking Databases (See also Surveillance)
- Remote Sensing
- Scene or Place Segmentation or Classification
- Segmentation
- Simultaneous Localization and Mapping
- Surveillance (See also People)
- Textures
- Urban Datasets
- Other Collection Pages
- Miscellaneous Topics

>600 Datasets Available at:

<http://homepages.inf.ed.ac.uk/rbf/CVonline/Imagedbase.htm>

FYI: CVonline

(<http://homepages.inf.ed.ac.uk/rbf/CVonline>)

has a lot of other resources:

- indexed pointers to about 2000 computer vision / image analysis topics
- list of image processing software
- list of image analysis applications
- list of about 950 computer vision / image processing / optics books
 - with links to 43 free online and 353 online subscription books

Project ideas

<http://cs231n.stanford.edu/project.html>

- [ICLR](#): International Conference on Learning Representations
- [Past CS229 Projects](#): Example projects from Stanford machine learning class
- [Kaggle challenges](#): An online machine learning competition website. For example, a [Yelp classification challenge](#).

For applications, this type of projects would involve careful data preparation, an appropriate loss function, details of training and cross-validation and good test set evaluations and model comparisons. Don't be afraid to think outside of the box. Some successful examples can be found below:

- [Teaching Deep Convolutional Neural Networks to Play Go](#)
- [Playing Atari with Deep Reinforcement Learning](#)
- [Winning the Galaxy Challenge with convnets](#)

ConvNets also run in real time on mobile phones and Raspberry Pi's - feel free to go the embedded way. You may find [DeepBeliefSDK](#) helpful. This particular project might be slightly out of date, but it may help you find more like it.

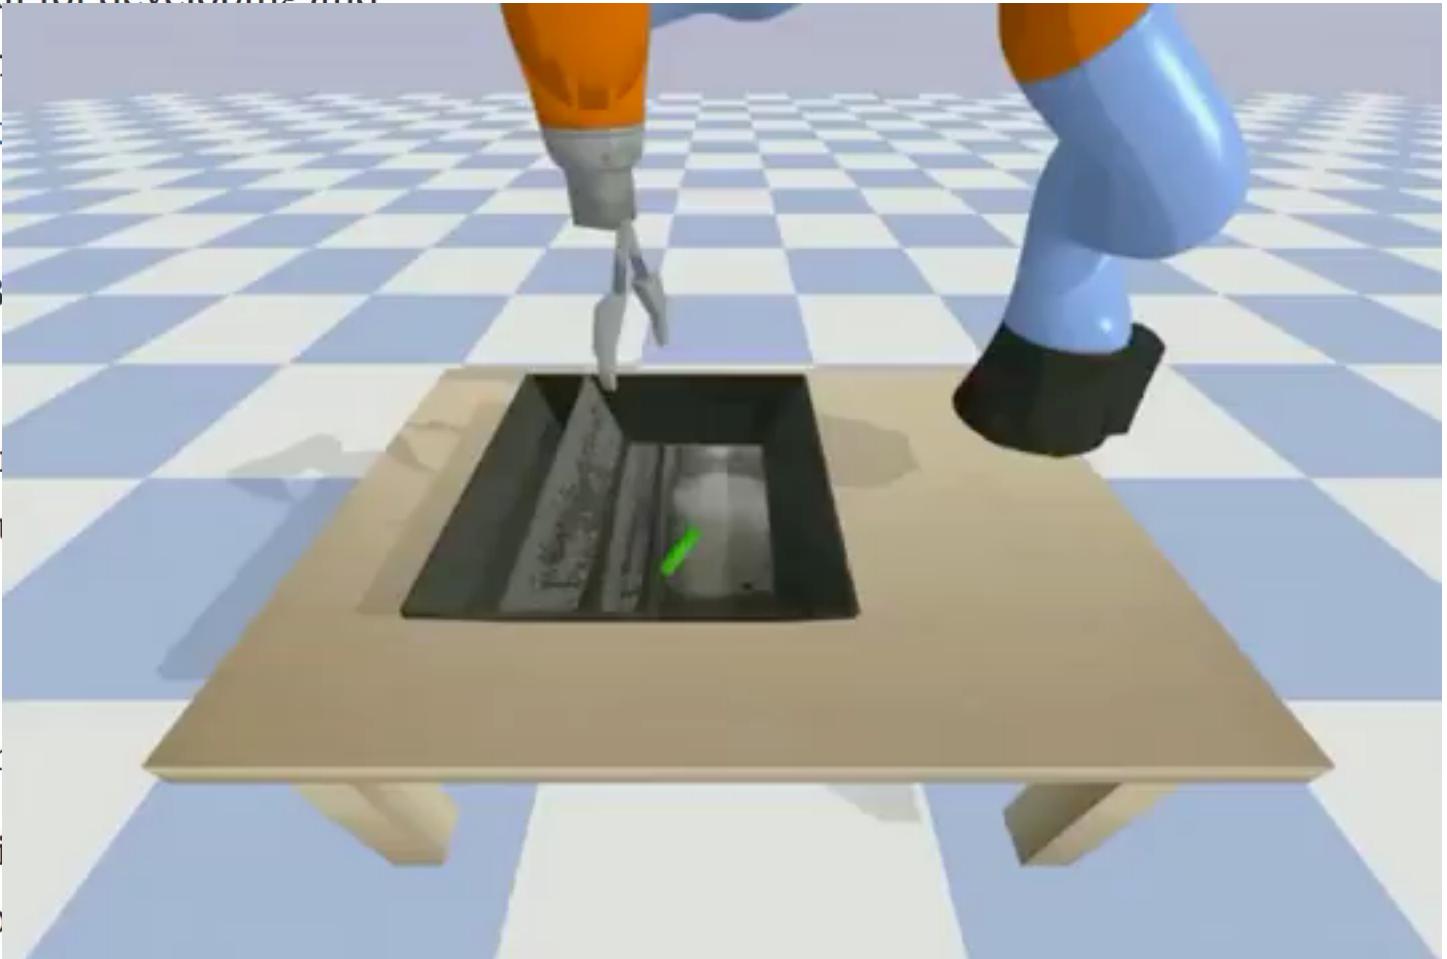
For models, ConvNets have been successfully used in a variety of computer vision tasks. This type of projects would involve understanding the state-of-the-art vision models, and building new models or improving existing models for a vision task. The list below presents some papers on recent advances of ConvNets in the computer vision community.

- **Object recognition**: [\[Krizhevsky et al.\]](#), [\[Russakovsky et al.\]](#), [\[Szegedy et al.\]](#), [\[Simonyan et al.\]](#), [\[He et al.\]](#)
- **Object detection**: [\[Girshick et al.\]](#), [\[Sermanet et al.\]](#), [\[Erhan et al.\]](#)
- **Image segmentation**: [\[Long et al.\]](#)
- **Video classification**: [\[Karpathy et al.\]](#), [\[Simonyan and Zisserman\]](#)
- **Scene classification**: [\[Zhou et al.\]](#)
- **Face recognition**: [\[Taigman et al.\]](#)
- **Depth estimation**: [\[Eigen et al.\]](#)
- **Image-to-sentence generation**: [\[Karpathy and Fei-Fei\]](#), [\[Donahue et al.\]](#), [\[Vinyals et al.\]](#)
- **Visualization and optimization**: [\[Szegedy et al.\]](#), [\[Nguyen et al.\]](#), [\[Zeiler and Fergus\]](#), [\[Goodfellow et al.\]](#), [\[Schaul et al.\]](#)

Project ideas

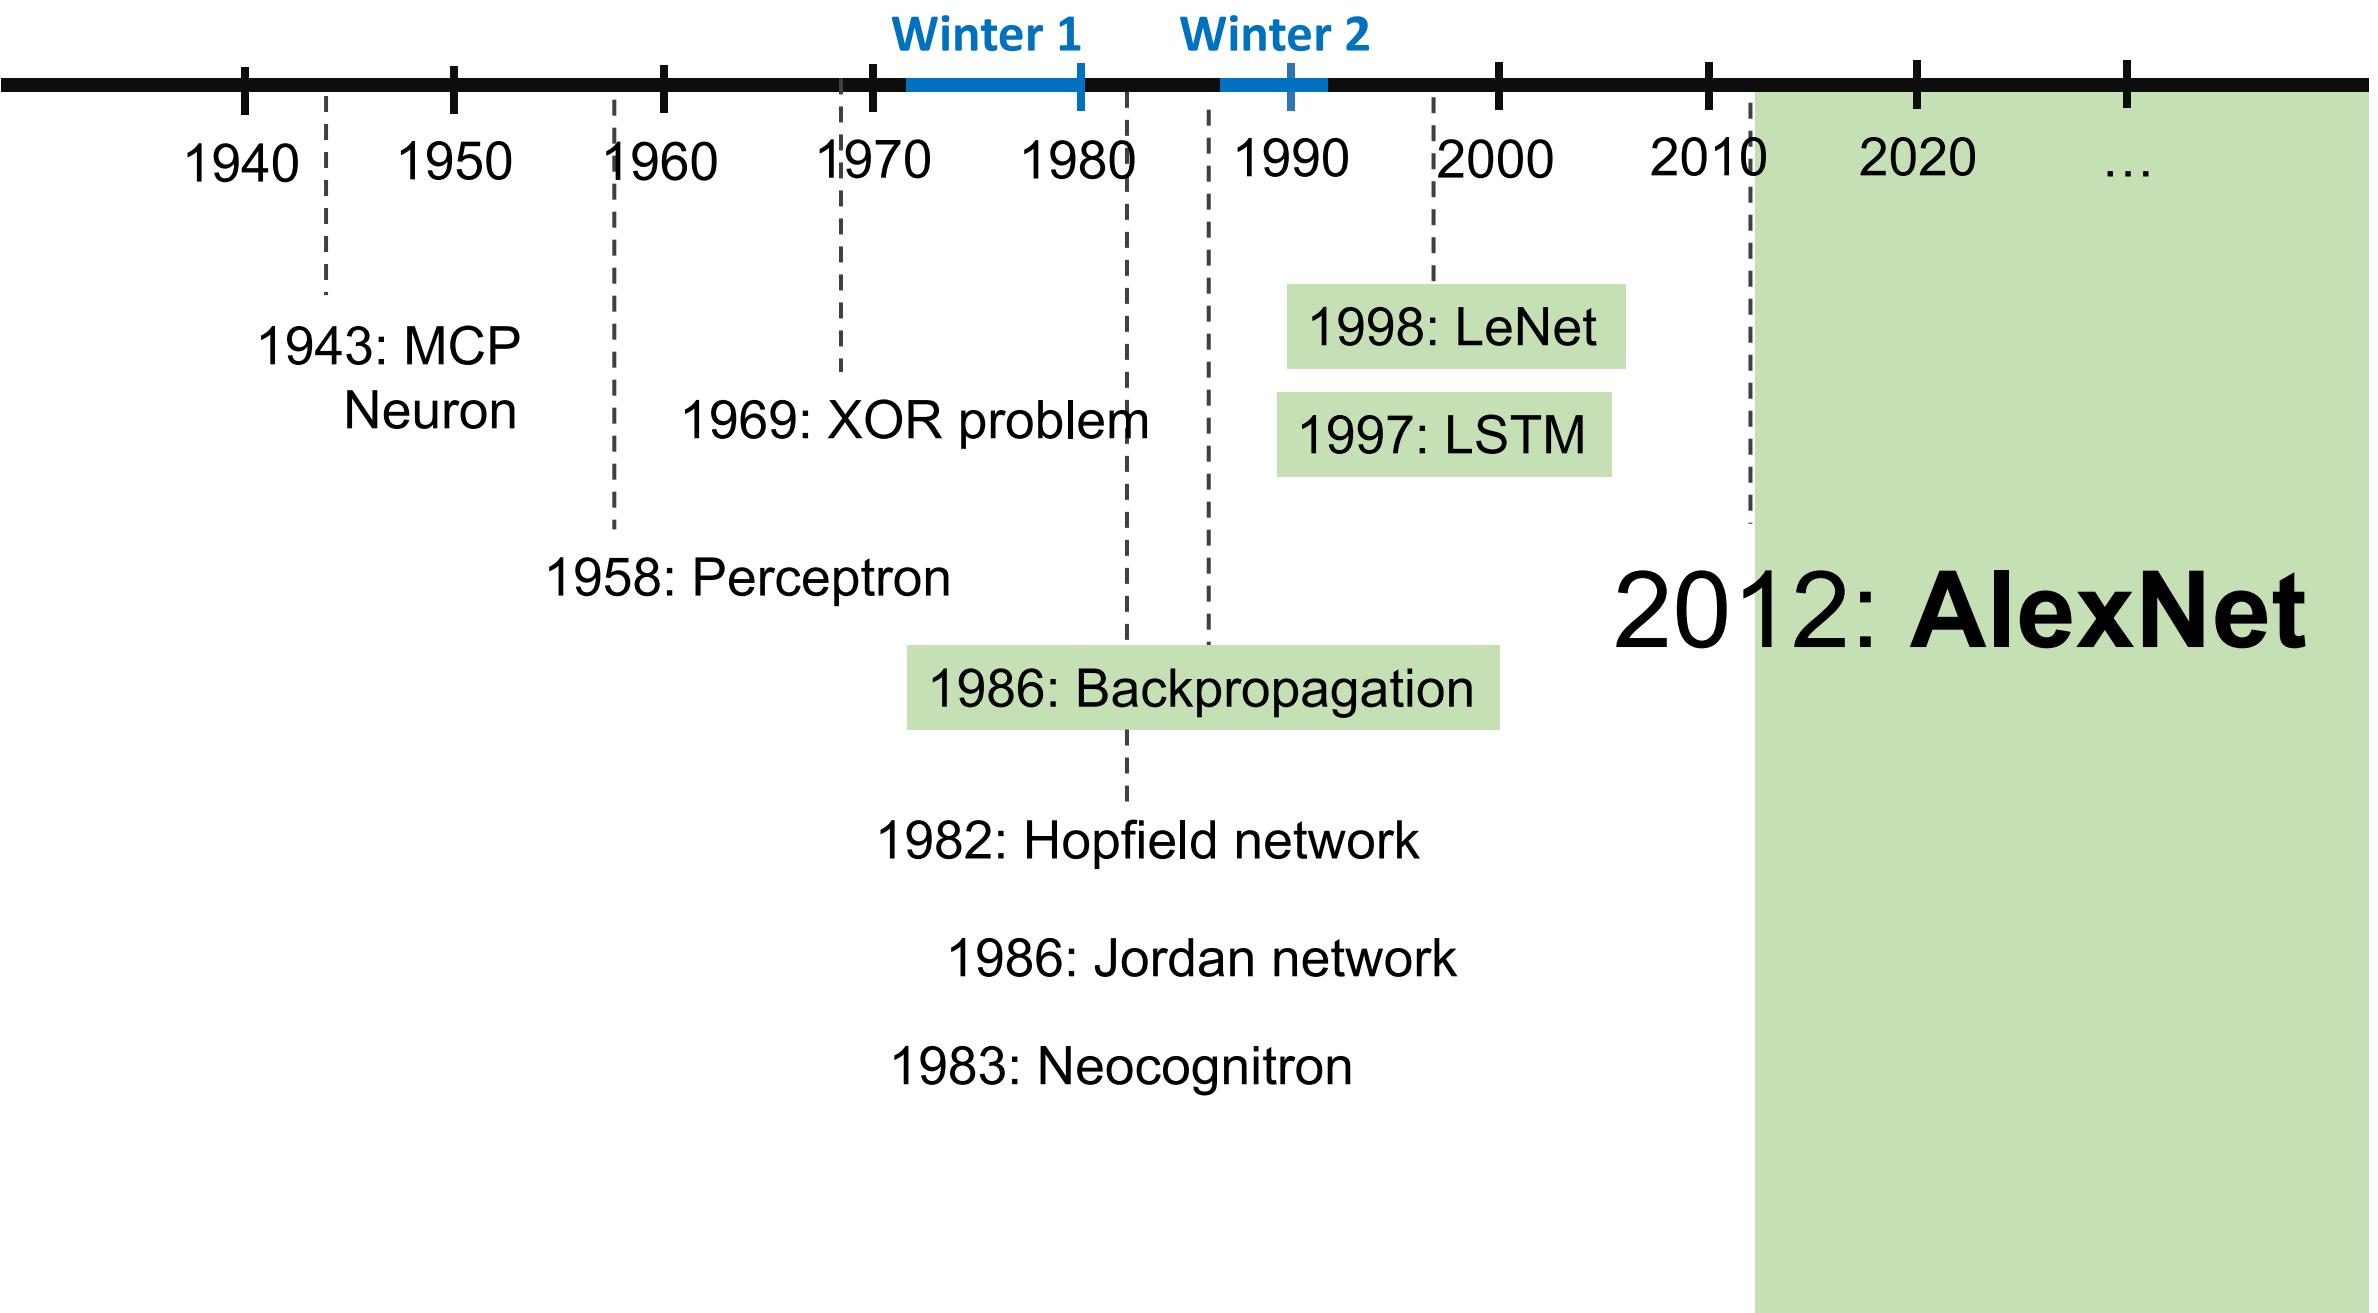
<https://www.quora.com/What-is-the-best-platform-to-simulate-Reinforcement-Learning-agent>

- OpenAI Gym [🔗](#) has become a standard toolkit for developing and comparing reinforcement learning algorithms across different environments. Now, it also integrates robotics simulation.
- DeepMind Lab [🔗](#) has a suite of challenging 3D environments for solving tasks for learning agents.
- OpenAI Universe [🔗](#) is a software platform for measuring AI's general intelligence across the world's scientific literature and other applications.
- DeepMind and Blizzard's Starcraft [🔗](#), which includes several minigames for developing AI agents that can play Starcraft II.
- The Arcade Learning Environment [🔗](#) (ALE) is a framework that allows researchers and hobbyists to train AI agents for Atari 2600 games.
- Microsoft's project Malmö [🔗](#) is a platform for Artificial Intelligence experimentation and research built on top of Minecraft. We aim to



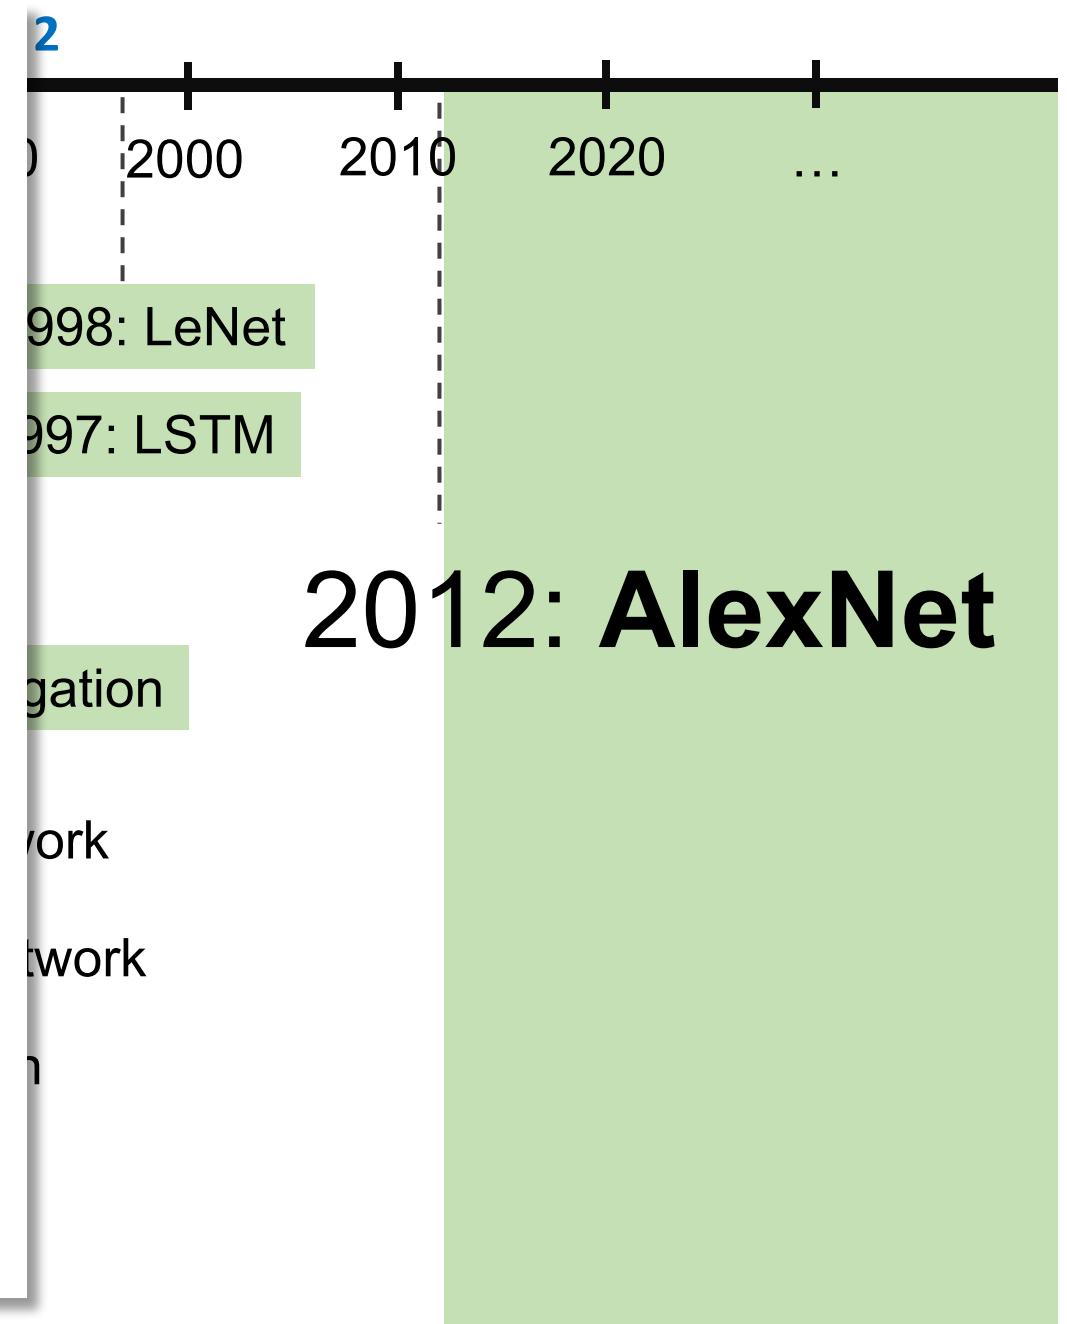
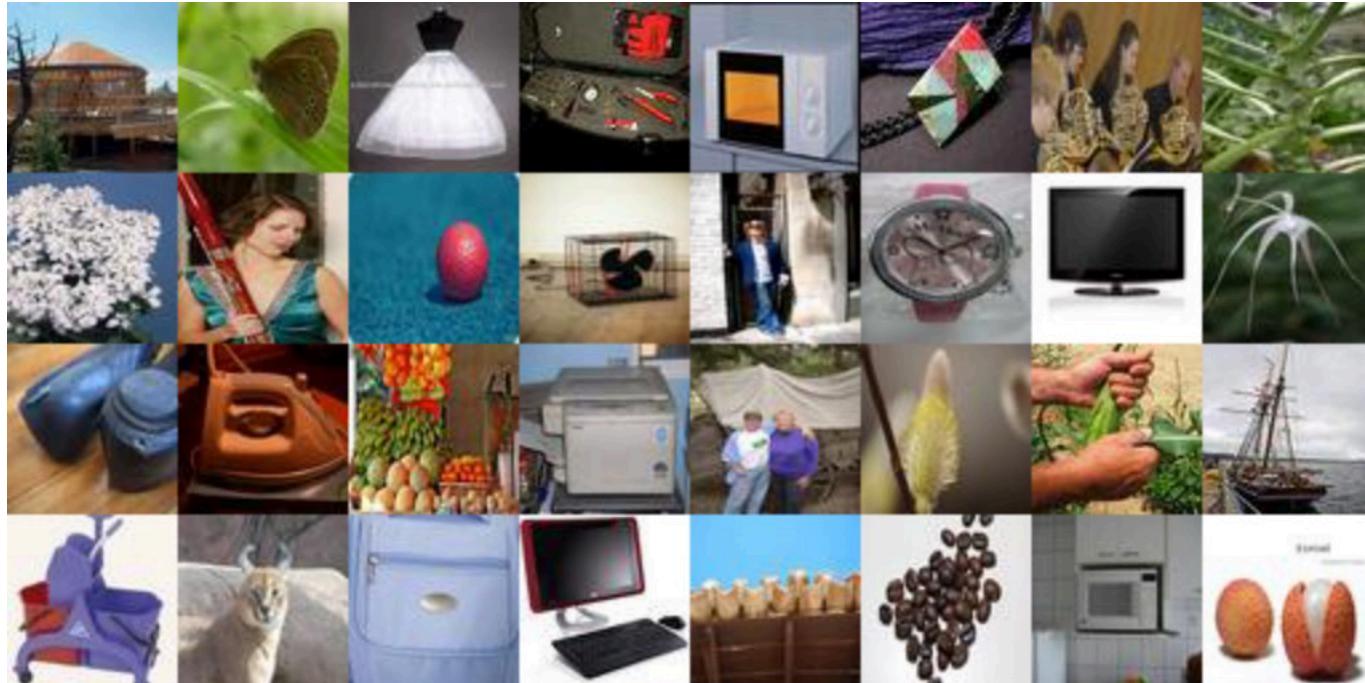
credit: David Ha

Anh Nguyen



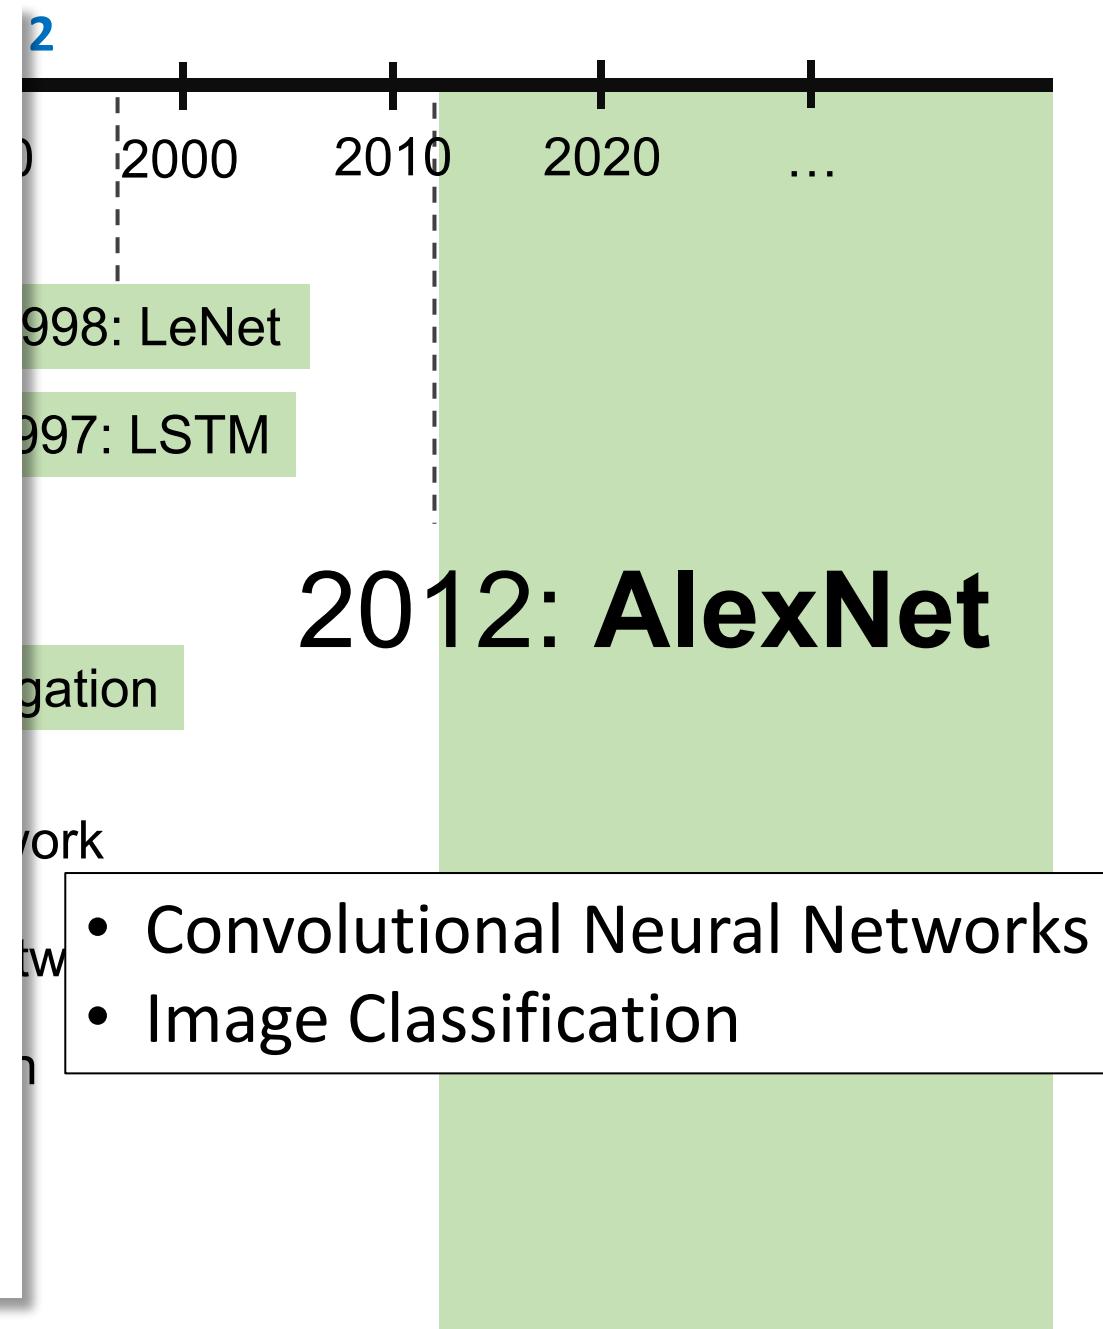
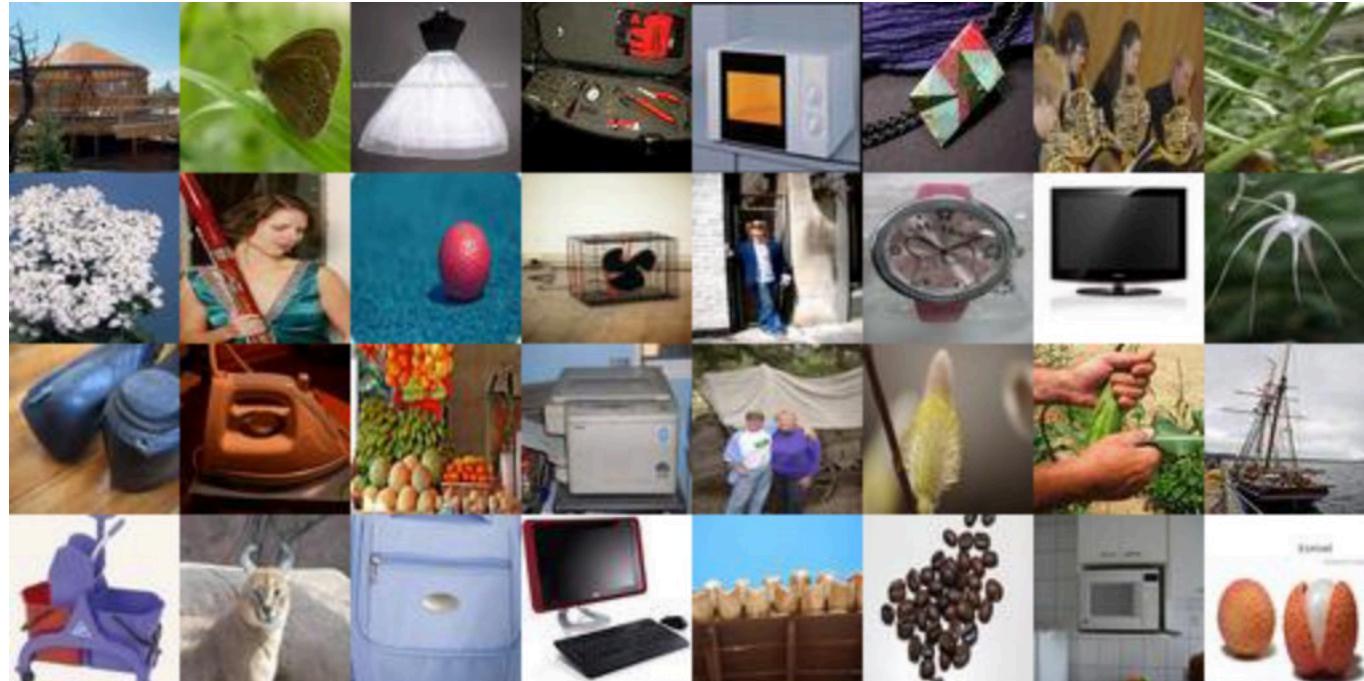
IMAGENET Challenge

- 1.3M images
- 1000 categories



IMAGENET Challenge

- 1.3M images
- 1000 categories



- 1.3M images
- 1000 categories



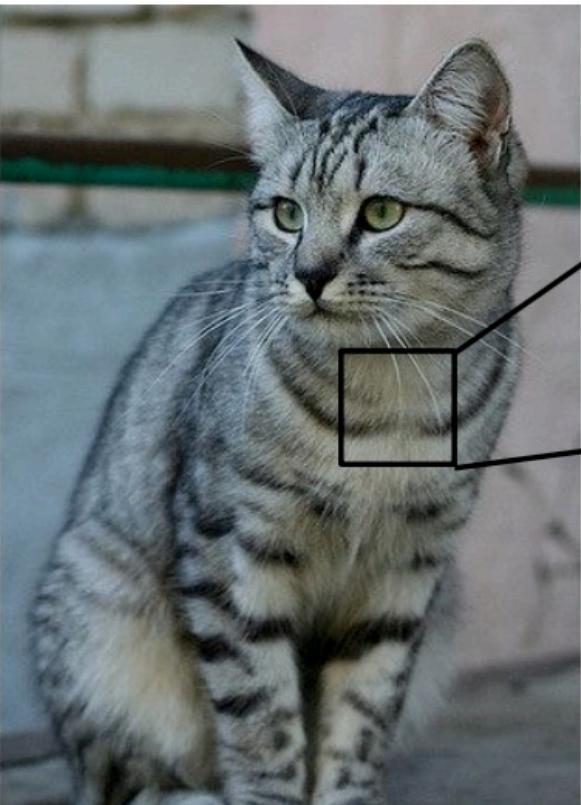
Image Classification

given set of discrete labels



dog
cat
truck
plane
...

The Problem: Semantic gap



This image by Nikita is
licensed under CC-BY 2.0

[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
[91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
[76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
[99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
[106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
[133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
[128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
[115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
[89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
[63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
[62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
[63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
[87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
[164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
[157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]
[130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
[128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
[123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
[122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]

What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Challenges: Illumination

credit: cs231n



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



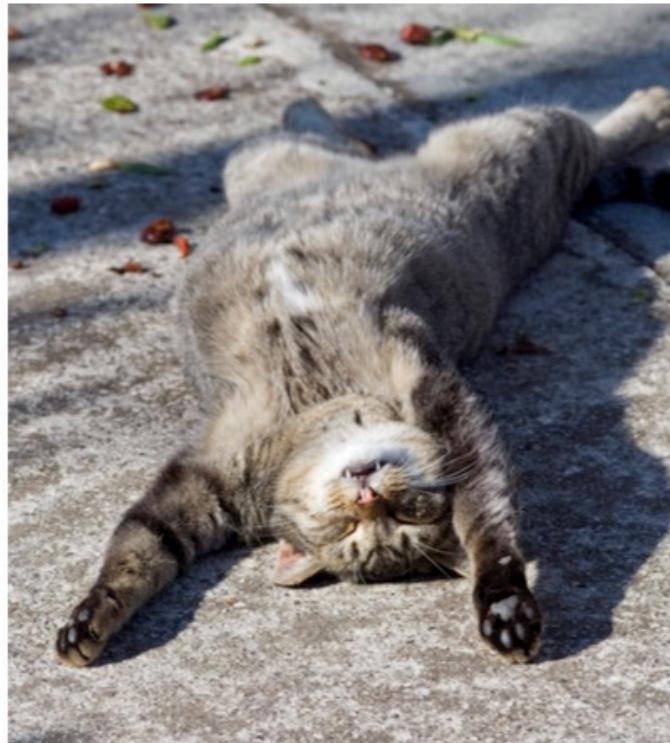
[This image is CC0 1.0 public domain](#)

Challenges: Deformation

credit: cs231n



[This image](#) by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



[This image](#) by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



[This image](#) by [sare bear](#) is
licensed under [CC-BY 2.0](#)

Challenges: Occlusion

credit: cs231n



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image by jonsson is licensed under CC-BY 2.0](#)

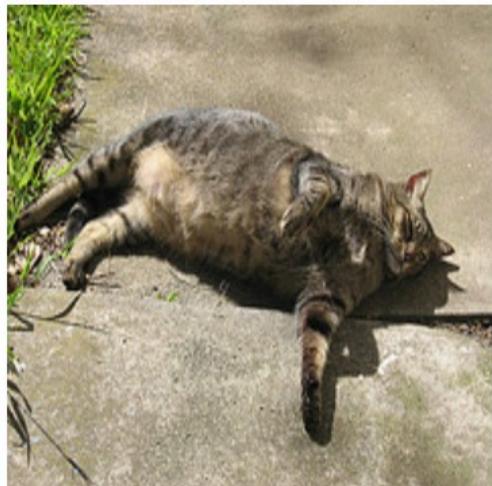
Challenges: Intra-class variation

credit: reddit



Challenges: Inter-class similarity

credit: ImageNet



Tabby cat

Tiger cat

An image classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for
recognizing a cat, or other classes.

Data-driven

1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images

Example training set

```
def train(images, labels):  
    # Machine learning!  
    return model
```

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```

airplane



automobile



bird



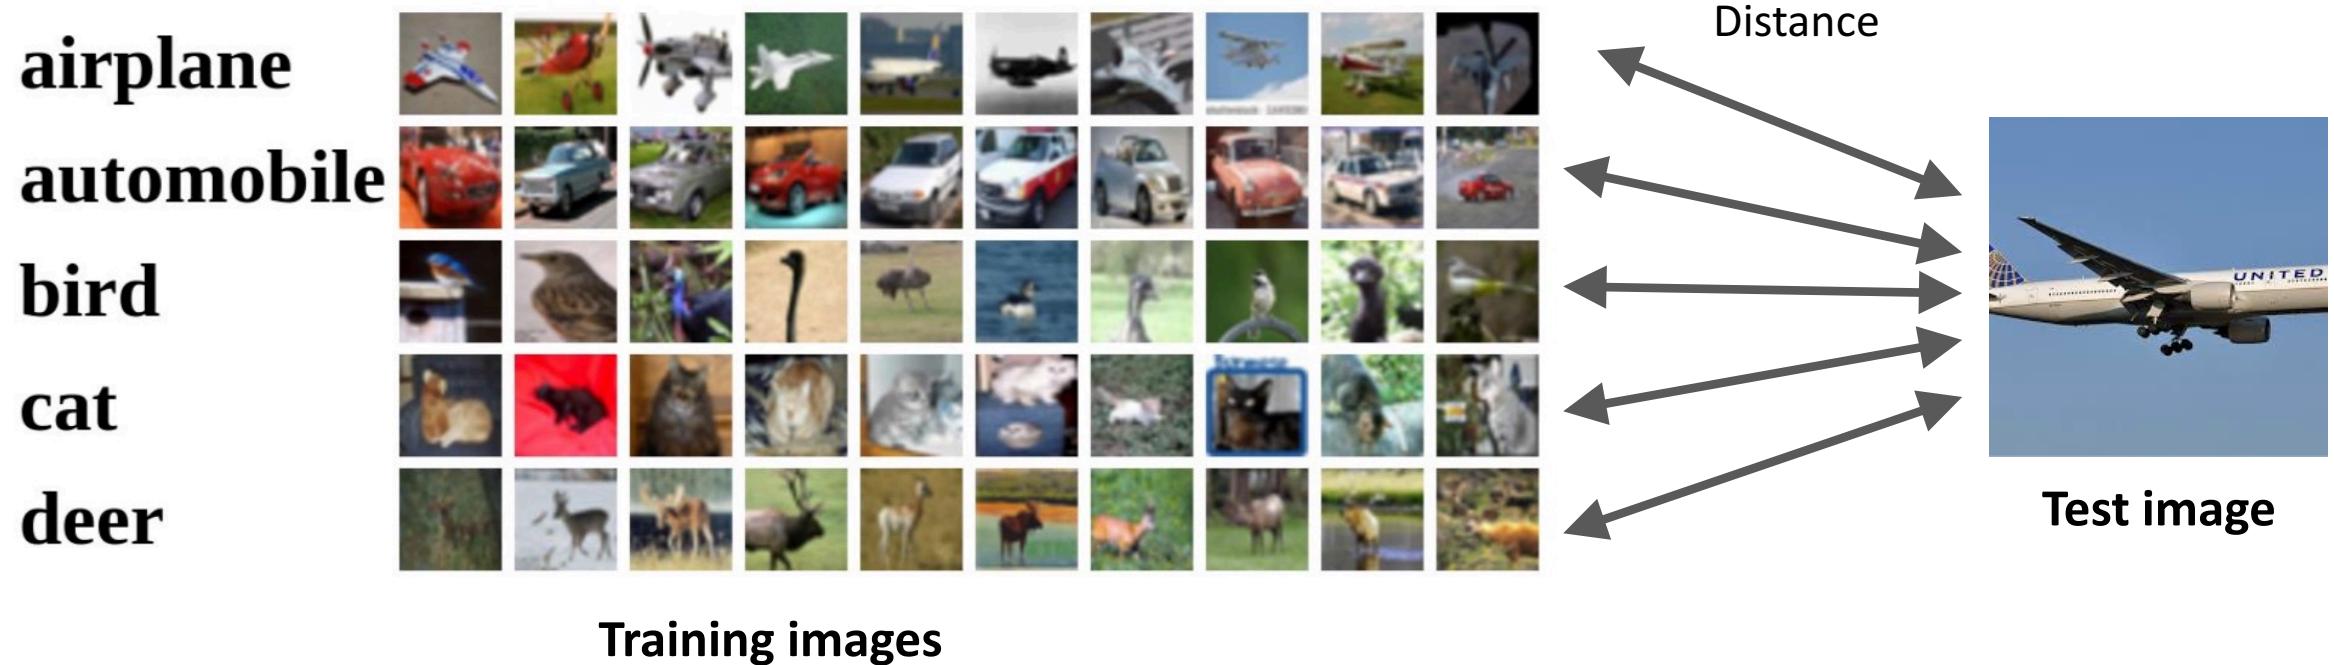
cat



deer



Nonparametric approach: Nearest Neighbor



Nonparametric approach: Nearest Neighbor

airplane



L1 distance:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image	56	32	10	18
-	90	23	128	133
training image	10	20	24	17
-	8	10	89	100
24	26	178	200	
-	12	16	178	170
2	0	255	220	
-	4	32	233	112

pixel-wise absolute value differences

46	12	14	1
82	13	39	33
12	10	0	30
2	32	22	108

add → 456

Distance



Test image

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

    return Ypred
```

Nearest Neighbor classifier

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

Memorize training data

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

For each test image:
 Find closest train image
 Predict label of nearest image

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

    return Ypred

```

Nearest Neighbor classifier

- Model size = Training set size
 - (*nonparametric*)
- Time complexity
 - Train: O(1)
 - Test: O(n)

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

- Model size = Training set size
 - (*nonparametric*)
- Time complexity
 - Train: O(1)
 - Test: O(n)
- Q: Is this a good classifier?

Demo:

<http://vision.stanford.edu/teaching/cs231n-demos/knn/>

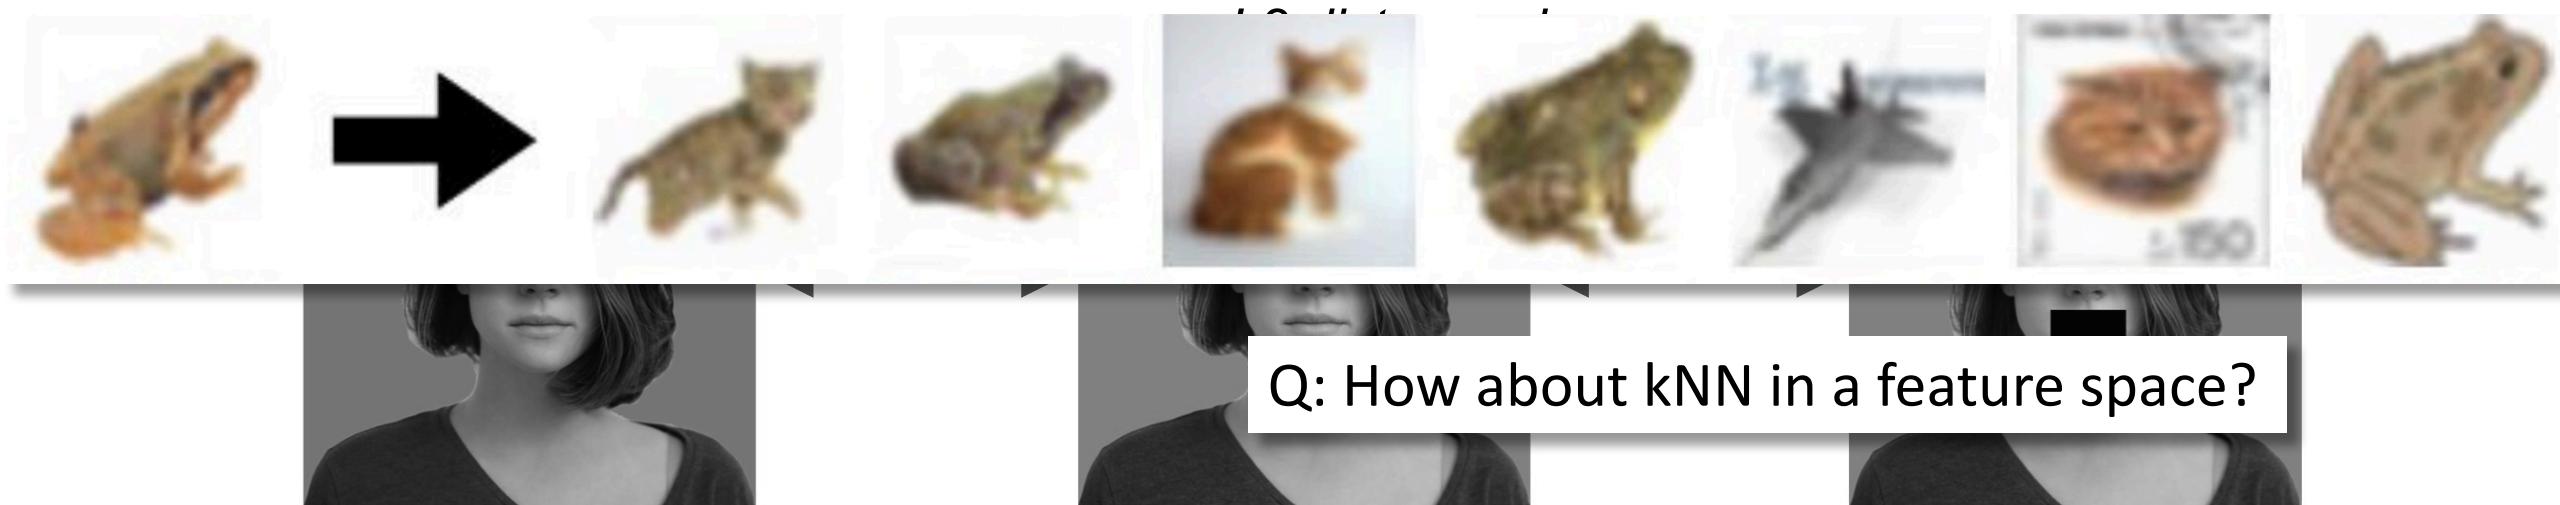
Problems with Nearest Neighbor

1. Model size (= training set size)
2. Slow test time
3. Curse of dimensionality
 - # training points needed is exponential wrt dimension
4. Distance metrics (e.g. L1, L2) in pixel space



Problems with Nearest Neighbor

1. Model size (= training set size)
2. Slow test time
3. Curse of dimensionality
 - # training points needed is exponential wrt dimension
4. Distance metrics (e.g. L1, L2) in pixel space

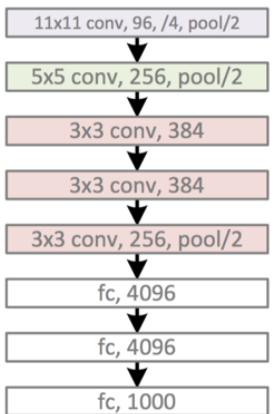


Linear Classifiers



2012: AlexNet

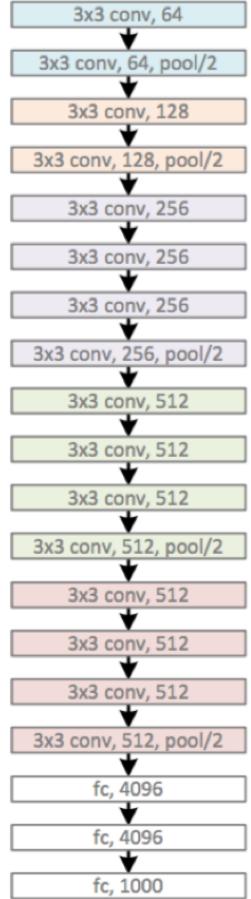
8 layers



Error: 15.3%

2014: VGG

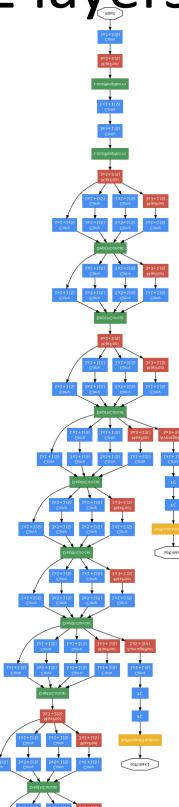
19 layers



8.5%

2015: GoogLeNet

22 layers



7.8%

2016: ResNet
>100 layers

2017

• • •

Human: 5.1%

4.4%

Neural Network

Linear
classifiers



Recall CIFAR10

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



50,000 training images
each image is **32x32x3**

10,000 test images.

Parametric Approach

Image



$$\xrightarrow{f(\mathbf{x}, \mathbf{W})}$$

\mathbf{W}
parameters
or weights

Array of **32x32x3** numbers
(3072 numbers total)

10 numbers giving
class scores

Parametric Approach: Linear Classifier

Image



$$f(x, W) = Wx$$

Array of **32x32x3** numbers
(3072 numbers total)

$$f(\mathbf{x}, \mathbf{W})$$

W
parameters
or weights

10 numbers giving
class scores

Parametric Approach: Linear Classifier

Image



Array of **32x32x3** numbers
(3072 numbers total)

$$f(x, W) = Wx$$

10x1 **10x3072**

3072x1

W
parameters
or weights

10 numbers giving
class scores

Parametric Approach: Linear Classifier

Image



Array of **32x32x3** numbers
(3072 numbers total)

$$f(x, W) = Wx + b$$

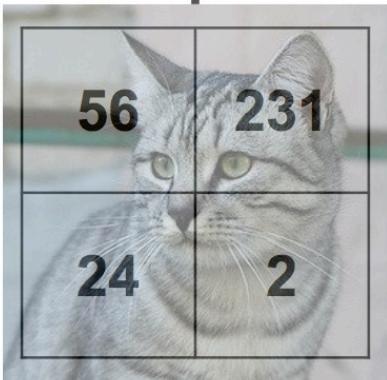
10x1 **3072x1**
10x3072 **10x1**

$f(x, W)$ → **10 numbers giving class scores**

W
parameters
or weights

Example with an image with 4 pixels, and 3 classes (cat/dog/**ship**)

Stretch pixels into column



Input image

0.2	-0.5	0.1	2.0
1.5	1.3	2.1	0.0
0	0.25	0.2	-0.3

W



+

1.1
3.2
-1.2

=

-96.8
437.9
61.95

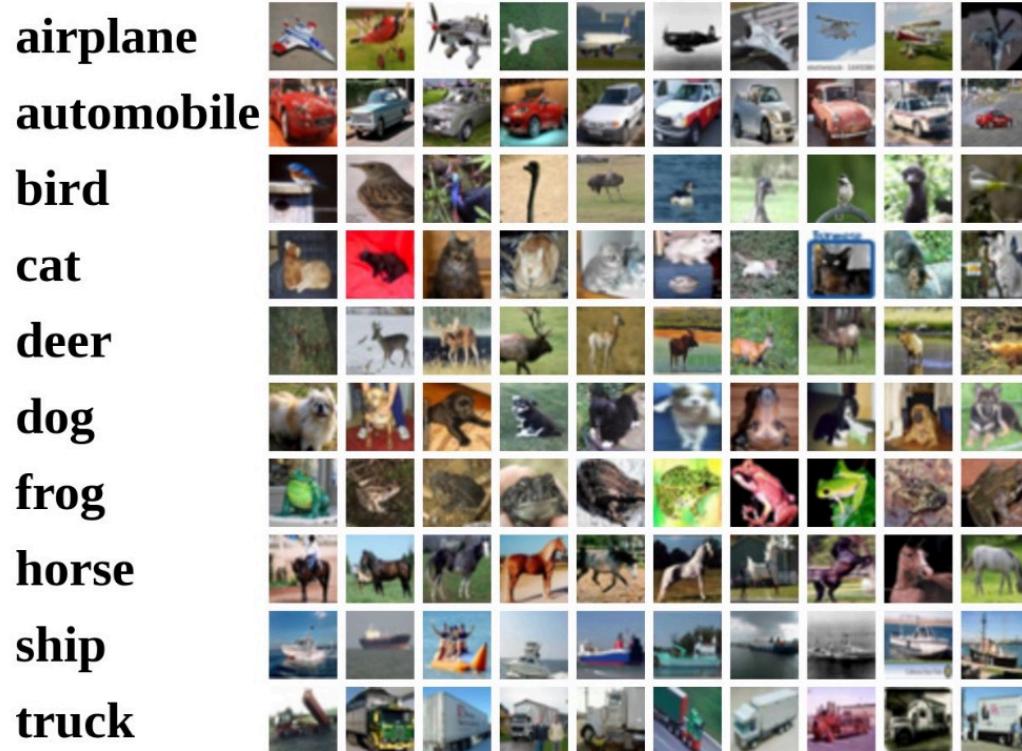
Cat score

Dog score

Ship score

Interpreting a Linear Classifier

Q: How to learn W?



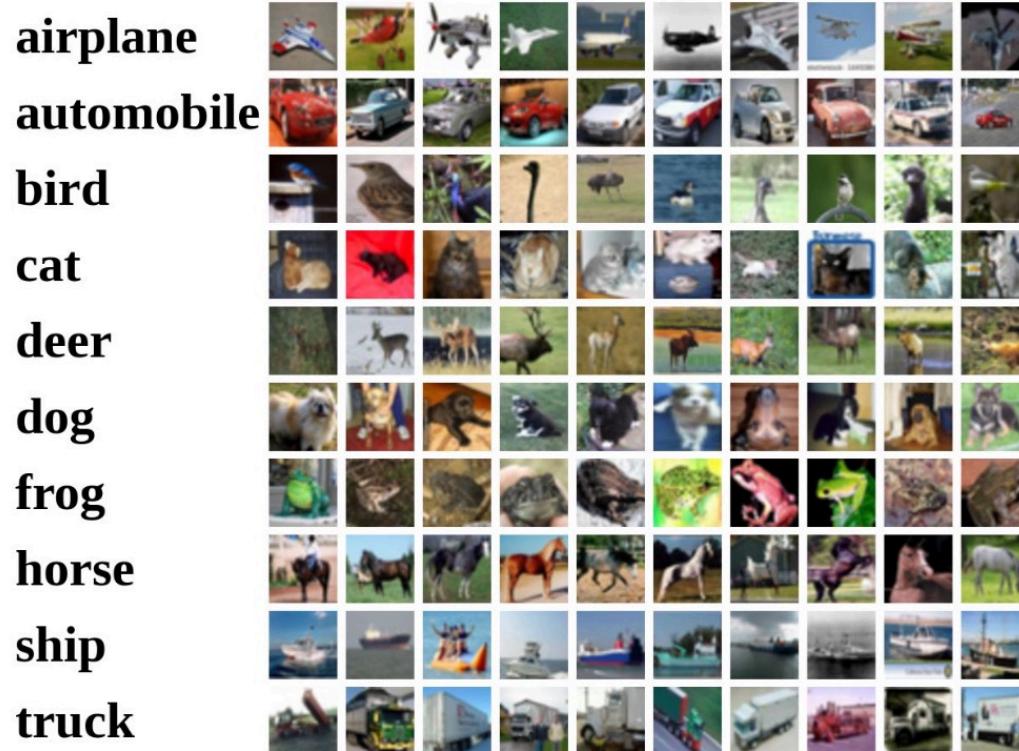
$$f(x, W) = Wx + b$$

Example trained weights
of a linear classifier
trained on CIFAR-10:



Interpreting a Linear Classifier

Q: How to learn W?



$$f(x, W)$$

Demo: <https://goo.gl/Qgr84h>
OR <http://playground.tensorflow.org/>

Example trained weights
of a linear classifier
trained on CIFAR-10:

